

Proofs

H.J. Rivera Verduzco 0977393

P.O. Box 513, 5600 MB Eindhoven, The Netherlands

Email: H.J.Rivera.Verduzco@student.tue.nl

May 3, 2017

1 Blocking tasks

Lemma 1. *Given a task-set \mathcal{T} of n independent tasks scheduled using FPTS, and a level- n busy period $[t_s, t_f]$ of such a schedule. If for a new task-set $\mathcal{T} \cup \{\tau_{n+1}\}$, τ_{n+1} is scheduled in such a way that a job k of τ_{n+1} executes within an interval $[t_1, t_s)$ where there is exactly an amount of time C_{n+1} available for the execution of τ_{n+1} , then the schedule of all jobs in $[t_s, t_f)$ remains the same after introducing τ_{n+1} .*

Proof. Assume that task τ_{n+1} is scheduled in such a way that a job k of τ_{n+1} executes within an interval $[t_1, t_s)$ where there is exactly an amount of time C_{n+1} available for the execution of τ_{n+1} . After introducing τ_{n+1} , the schedule in $[t_s, t_f)$ can be guaranteed to remain unchanged if there is no pending load at time t_s from higher priority tasks of τ_{n+1} , i.e. if $P_n(t_s) = 0$. In addition, there should not be pending load at time t_s from a job of τ_{n+1} that has already started. We show that these two conditions are always met for the given assumptions.

$\{P_n(t_s) = 0\}$. First note that, since there is exactly an amount of time C_{n+1} available for the execution of τ_{n+1} in the interval $[t_1, t_s)$, job k of τ_{n+1} can fit exactly in such an interval. Therefore, the higher priority tasks executing in $[t_1, t_s)$ will still finalize before or at t_s , and consequently the pending load $P_n(t_s)$ will be zero.

$\{\text{Pending load of already started jobs of } \tau_{n+1} \text{ at time } t_s \text{ is zero}\}$. Since we assumed that job k executes within $[t_1, t_s)$ and therefore it finalizes before or at t_s , we only have to prove that next jobs of τ_{n+1} cannot start before t_s . Note that job k of τ_{n+1} already occupies all the time that was idle in $[t_1, t_s)$. Therefore, the start time of any job of τ_{n+1} after job k activated before t_s will be postponed after t_f . □

Lemma 2. *Given a task-set \mathcal{T} of n independent tasks scheduled under FPTS, a level- n busy period $[t_s, t_f]$ of such a schedule, and a new task-set $\mathcal{T} \cup \{\tau_{n+1}\}$ where $U^{T \cup \tau_{n+1}} < 1$ or $U^{T \cup \tau_{n+1}} = 1$ and the lcm of the periods exists. It is always possible for task-set $\mathcal{T} \cup \{\tau_{n+1}\}$ to schedule τ_{n+1} in such a way that a job k of τ_{n+1} executes within an interval $[t_1, t_s)$ where there is exactly an amount of time C_{n+1} available for the execution of τ_{n+1} .*

Proof. Since $U^{T \cup \tau_{n+1}} < 1$ or $U^{T \cup \tau_{n+1}} = 1$ and the lcm of the periods exists, it is always possible to find a contingent interval $[t_1, t_s)$ before the level- n busy period $[t_s, t_f]$ where there is an amount of time C_{n+1} available for the execution of τ_{n+1} . The only reason for which a job k of τ_{n+1} may not be able to execute within $[t_1, t_s)$ is that it could experience some blocking from previous jobs delaying the start time of such a job k . However, it is possible to change the phase of ϕ_{n+1} in order to move the activations of the jobs of τ_{n+1} at an earlier moment in time, and therefore allow enough space before $[t_1, t_s)$ for the execution of the earlier jobs of τ_{n+1} before job k that may delay its start time. Since the number of jobs of τ_{n+1} that may delay the start time of job k is bounded by the number of jobs in the worst-case length of the level- $(n+1)$ active period, it is always possible to find enough space before $[t_1, t_s)$ for their execution. □

Lemma 3. *Let all tasks of a set \mathcal{T} be strictly periodic and let $U^T < 1$ or $U^T = 1$ and the lcm of the periods exists. Under FPTS, the best-case response time BR_i of a task $\tau_i \in \mathcal{T}$ is not influenced by its lower priority tasks.*

Proof. Assume that we find a schedule for a subset of task-set \mathcal{T} defined as $\mathcal{T}_i = \{\tau_a | a \leq i, \tau_a \in \mathcal{T}\}$ where the best-case response time BR_i of τ_i is assumed by job k^{bcr} in the level- i busy period $[t_s, t_f]$. Note that task-set \mathcal{T}_i contains the same tasks as \mathcal{T} without the lower priority tasks of τ_i . Given lemmas 1 and 2, we can construct a schedule for a task-set \mathcal{T}_{i+1} where task τ_{i+1} is scheduled in such a way that a job k executes in a contingent time interval $[t_1, t_s)$ before $[t_s, t_f]$. Hence, the schedule in time interval $[t_s, t_f]$ will remain unchanged and the best-case response time BR_i of τ_i will remain the same in \mathcal{T}_{i+1} . Furthermore, note that since job k of τ_{i+1} occurs in a contingent time interval $[t_1, t_s)$ before $[t_s, t_f]$, these two time intervals are contained in a level- $\{i+1\}$ busy period $[t_s, t'_f]$; hence, job k^{bcr} of τ_i occurs in $[t'_s, t'_f]$ as well. Based on this observation and using lemmas 1 and 2 again, we can find a schedule for a task-set \mathcal{T}_{i+2} where the jobs in the level- $\{i+1\}$ busy period $[t'_s, t'_f]$ remain unchanged and, therefore, BR_i is preserved. Since we can continue constructing schedules preserving the best-case response time of τ_i until $\mathcal{T}_{i+l} = \mathcal{T}$, we conclude that the best-case response time BR_i of a task τ_i is not affected by its lower priority tasks. \square

2 Higher priority tasks

Lemma 4. *Let \mathcal{T} be a set of strictly periodic tasks with a task $\tau_i \in \mathcal{T}$. The threshold of higher priority tasks do not influence the response time of a job of τ_i .*

Proof. In order to prove Lemma 4, first note that the response time $R_{i,k}$ of a job $\iota_{i,k}$ of τ_i is equal to $S_{i,k} + H_{i,k}$. Therefore, it is sufficient to show that the thresholds of higher priority tasks do not have an influence on the start time $S_{i,k}$ and the hold time $H_{i,k}$.

{Influence on $S_{i,k}$ }. Under FPTS, a necessary condition for a job $\iota_{i,k}$ to start at a time t after its activation is that the priority of τ_i is the highest among all tasks with pending load at time t . Since this condition only depends on the priority of the tasks, higher priority tasks will always affect the start time $S_{i,k}$ of $\iota_{i,k}$ independently of their thresholds.

{Influence on $H_{i,k}$ }. Under FPTS, a job $\iota_{i,k}$ of a task τ_i can only be preempted by a higher priority task τ_h iff $\pi_h > \theta_i$. Since this condition is independent of the threshold θ_h of τ_h , we conclude that the thresholds of higher priority tasks of τ_i do not have an influence on the hold time $H_{i,k}$ of $\iota_{i,k}$. \square

3 BCRT for non-preemptive tasks

Lemma 5. *Let τ_i be a non-preemptive task. The best-case response time of τ_i is always equal to its best-case computation time, i.e. $BR_i = BC_i$.*

Proof. Given a task-set \mathcal{T} and a non-preemptive task $\tau_i \in \mathcal{T}$, assume that all jobs of τ_i have a computation time equal to BC_i . Since τ_i is non-preemptive, the hold time of all its jobs is simply equal to BC_i . Recall that the response time of a job k of τ_i is given by $R_{i,k} = S_{i,k} + H_{i,k}$; therefore, in order to prove the lemma, it is sufficient to prove that it is always possible to schedule the tasks in such a way that the relative start time of a job k of τ_i is zero, i.e. $S_{i,k} = 0$. We divide this proof in two cases:

{Case $U^T < 1$ }. Given an arbitrary schedule for \mathcal{T} , let $[t_s, t_e)$ be an idle interval in such a schedule. Note that it is always possible to find an idle interval because $U^T < 1$. Furthermore, let $\iota_{i,k}$ be the first job of τ_i activated at or after time t_e , i.e. $t_e \leq a_{i,k}$ and there are no jobs of τ_i activated in $[t_e, a_{i,k})$. We now show that by pushing the activation of $\iota_{i,k}$ to occur in the interval $[t_s, t_e)$, the relative start time $S_{i,k}$ of job $\iota_{i,k}$ becomes zero.

First observe that there is no other job of τ_i between the idle interval $[t_s, t_e)$ and $\iota_{i,k}$; therefore, after pushing the activations of τ_i till $a_{i,k}$ occurs in $[t_s, t_e)$, job $\iota_{i,k}$ cannot experience blocking by a previous job. Furthermore, since $[t_s, t_e)$ was originally idle, job $\iota_{i,k}$ can start upon activation, leading to $S_{i,k} = 0$.

{Case $U^T = 1$ and the lcm of the periods exists}. Given an arbitrary schedule for \mathcal{T} , we first observe that the schedule will be repeated every hyper-period because the lcm of the period exists. Assume that there are n jobs of τ_i in every hyper-period and no job of τ_i starts upon activation. Furthermore, let $\iota_{i,1}$ and $\iota_{i,n}$ be the first and last job of τ_i respectively in an hyper-period. Hence, for each interval $[a_{i,k'}, s_{i,k'})$ with $1 \leq k' \leq n$, there are only higher priority tasks of τ_i and/or at most one blocking lower priority task executing at that interval. Let $\iota_{i,k}$ be the job of τ_i with the shortest relative start time among all jobs of τ_i in an hyper-period. Since no job of τ_i starts upon activation and $\iota_{i,k}$ is the job with the shortest start time, we can push the activation of all jobs of τ_i to a later moment in time, without changing the schedule, till job $\iota_{i,k}$ starts upon activation. After pushing the activations of τ_i , it holds that $S_{i,k} = 0$ therefore concluding the proof. □

References