

# Towards best-case response times of real-time tasks under fixed-priority scheduling with preemption thresholds

H.J. Rivera Verduzco

*Technische Universiteit Eindhoven (TU/e),  
Den Dolech 2, 5600 AZ Eindhoven,  
The Netherlands  
h.j.rivera.verduzco@student.tue.nl*

Reinder J. Bril

*Technische Universiteit Eindhoven (TU/e),  
Den Dolech 2, 5600 AZ Eindhoven,  
The Netherlands  
r.j.bril@tue.nl*

**Abstract**—Fixed-priority scheduling with preemption thresholds (FPTS) is currently supported by AUTOSAR and OSEK standards as an scheduling policy. Since FPTS is a generalization for fixed-priority preemptive scheduling (FPPS) and for fixed-priority non-preemptive scheduling (FPNS), it aims to improve schedulability while reducing preemption overhead. In this paper, we show that the best-case response time analysis for FPTS is most likely not a straight forward extension of the current best-case analysis for FPPS. In addition, we show, as an intermediate step towards the exact best-case response time analysis for FPTS, that the execution time of a task scheduled under FPNS is a tight lower bound for its response time.

## 1. Introduction

Fixed-priority scheduling with preemption thresholds (FPTS) was first introduced in the ThreadX kernel [1] to limit the number of preemptions that a task may experience with the aim to improve schedulability and reduce preemption overheads. It is currently also supported by automotive standards AUTOSAR and OSEK [2, 3]. Although the worst-case response time analysis for FPTS has been studied extensively in the literature [4-6], little attention has been paid to its best-case counterpart.

FPTS is a generalization for fixed-priority preemptive scheduling (FPPS) and fixed-priority non-preemptive scheduling (FPNS). Therefore, it could be desirable to consider first the best-case response time analysis for FPPS and FPNS before investigating the corresponding analysis for FPTS. In particular, the best-case response time analysis for FPPS has been addressed in [7] for constraint deadlines and in [8] for arbitrary deadlines. On the other hand, the best-case response time for FPNS has not been formally studied.

FPNS is widely used in message-based protocols under a network environment. For example, the Controller Area Network (CAN) uses a priority based arbitration and non-preemptive scheduling for message transmission. Furthermore, special attention has been paid to the worst-case analysis for CAN in [9], and for FPNS in [10] to guarantee real-time requirements. However, when considering jitter induced by a network, it becomes relevant to study best-

case response times next to worst-case response times of messages scheduled under FPNS.

In [11], a lower bound for the best-case response time under fixed-priority scheduling with deferred preemptions (FPDS) is presented. Since FPNS is a special case of FPDS, this analysis also applies for FPNS. In particular, this analysis gives the trivial result that the execution time of a task is a lower bound for the best-case response time of such a task under FPNS. However, we would like to investigate whether this lower bound is tight for all cases.

The rest of this paper is organized as follows. Section 2 introduces the basic real-time scheduling model for FPTS that we will analyze as well as some related notions. Section 3 introduces an example that refutes an initial and intuitive notion of optimal instant for FPTS. The best-case response time for FPNS is presented in Section 4. Finally, we conclude this paper in Section 5.

## 2. Real-time scheduling model

In this section, we present the real-time scheduling model for FPTS along with some related notions.

### 2.1. Basic model for FPTS

We assume a single processor and a set  $\mathcal{T}$  of  $n$  independent periodic tasks  $\tau_1, \tau_2, \dots, \tau_n$  with unique and fixed priorities  $\pi_1, \pi_2, \dots, \pi_n$ . We also assume that tasks are given in order of decreasing priority, i.e.  $\tau_1$  has the highest priority whereas  $\tau_n$  has the lowest priority. A higher priority is represented by a higher value.

Each task  $\tau_i$  generates an infinite sequence of jobs  $\iota_{ik}$  with  $k \in \mathbb{Z}$ . In addition, each task is characterized by a *period*  $T_i \in \mathbb{R}^+$ , a *worst-case computation time*  $WC_i \in \mathbb{R}^+$ , a *best-case computation time*  $BC_i \in \mathbb{R}^+$ , where  $BC_i \leq WC_i$ , a *phasing*  $\phi_i \in \mathbb{R}$ , a (relative) *worst-case deadline*  $WD_i \in \mathbb{R}^+$ , and a (relative) *best-case deadline*  $BD_i \in \mathbb{R}^+ \cup \{0\}$ , where  $BD_i \leq WD_i$ . We assume arbitrary phasings. Moreover, the set of phasings  $\phi_i$  is termed the phasing  $\phi$  of the task-set  $\mathcal{T}$ . Deadlines  $BD_i$  and  $WD_i$  are relative to the activations of the jobs. Furthermore, we assume arbitrary deadlines; hence, deadline  $WD_i$  may be smaller than, equal to, or larger




Figure 1. Basic model for a periodic task  $\tau_i$ .

than period  $T_i$ . In addition, we assume that tasks do not suspend themselves, jobs do not start before the completion of previous jobs of the same task, and the overhead of context switching and task scheduling is ignored.

For fixed-priority preemptive scheduling with preemption thresholds (FPTS), each task  $\tau_i$  has an additional property called *preemption threshold* denoted by  $\theta_i$ , where  $\pi_1 \geq \theta_i \geq \pi_i$ . A task  $\tau_i$  can be preempted by a higher priority task  $\tau_h$  if and only if  $\pi_h > \theta_i$ . FPPS is a special case of FPTS when  $\forall_{1 \leq i \leq n} \theta_i = \pi_i$ , and FPNS is a special case when  $\forall_{1 \leq i \leq n} \theta_i = \pi_1$ .

It is worth noting that activation jitter is ignored in the real-time scheduling model under study. We aim to investigate first the best-case response times analysis for FPTS for the basic model before refining the model with the notion of activation jitter.

## 2.2. Derived notions

Given the *activation time*  $a_{i,k}$  of the  $k^{th}$  job of a task  $\tau_i$  and its (absolute) *finalization time*  $f_{i,k}$ , the *response time*  $R_{i,k}$  of such a job is defined as the time elapsed between its finalization and activation, i.e.  $R_{i,k} = f_{i,k} - a_{i,k}$ . The *absolute start time*  $s_{i,k}$  is the time at which job  $k$  of  $\tau_i$  starts its execution. Furthermore, the *relative start time*  $S_{i,k}$  is the start time relative to the activation of a job, i.e.  $S_{i,k} = s_{i,k} - a_{i,k}$ . Finally, we define the *hold time*  $H_{i,k}$  as the time elapsed between the start of a job and its completion, this is expressed as  $H_{i,k} = f_{i,k} - s_{i,k}$ . In general, under this model, it always holds that  $R_{i,k} = S_{i,k} + H_{i,k}$ . Figure 1 shows the basic task model with these notions.

The *worst-case response time*  $WR_i$  and the *best-case response time*  $BR_i$  of a task  $\tau_i$  are defined as the longest and the shortest response times of its jobs respectively, i.e.

$$WR_i \stackrel{\text{def}}{=} \sup_{\phi, k} R_{i,k}(\phi) \quad (1)$$

$$BR_i \stackrel{\text{def}}{=} \inf_{\phi, k} R_{i,k}(\phi), \quad (2)$$

where  $R_{i,k}(\phi)$  denotes a dependency of response time  $R_{i,k}$  on phasing  $\phi$ .

We say that a set  $\mathcal{T}$  of  $n$  periodic tasks is schedulable if and only if

$$\forall_{1 \leq i \leq n} (BD_i \leq BR_i \wedge WR_i \leq WD_i). \quad (3)$$

Finally, we introduce the notion of level- $i$  active period and best-case utilization. In order to define the notion of level- $i$  active period [10], we first introduce the notion of pending load. The *pending load*  $P_i(t)$  of a task  $\tau_i$  is defined as the amount of processing at time  $t$  that still needs to be completed for the jobs with a priority higher than or equal to task  $\tau_i$  that are activated before time  $t$ . A level- $i$  active period is then defined as the interval  $[t_s, t_e)$  such that  $P_i(t_s) = 0$ ,  $P_i(t_e) = 0$ , and  $P_i(t) > 0$  for all  $t \in (t_s, t_e)$ .

We define the best-case utilization  $U^T$  as the best-case fraction of the processor time spent on the execution of  $\mathcal{T}$ , i.e.

$$U^T \stackrel{\text{def}}{=} \sum_{1 \leq i \leq n} \frac{BC_i}{T_i}. \quad (4)$$

## 3. Towards best-case analysis for FPTS

In this section, we illustrate by means of an example that the best-case response time analysis for FPTS is most likely not a trivial extension of the existing best-case analysis for FPPS.

### 3.1. Recap on optimal instant for FPPS and FPNS

An optimal instant is defined as the (hypothetical) instant that leads to the best-case response time of a task. For FPPS, the notion of optimal instant is the following [8]:

**Theorem 1.** *An optimal instant of a task  $\tau_i$  under FPPS and arbitrary deadlines occurs when the completion of a job of task  $\tau_i$  coincides with the simultaneous activation of all its higher priority tasks.*

It is worth noting that, for FPPS, the job experiencing the optimal instant is the job that assumes the best-case response time.

For FPNS, a conjecture of a so called  $\Delta$ -optimal instant is presented in [11]. Based on this conjecture, a lower bound for the best-case response time under FPNS is also derived.

**Conjecture 1.** *A  $\Delta$ -optimal instant of a task  $\tau_i$  under FPNS occurs when a job of  $\tau_i$  starts a sufficiently small amount of time  $\Delta > 0$  before the simultaneous activation of all its higher priority tasks.*




Figure 2. A timeline for  $\mathcal{T}_1$  depicting a level-3 active period. The activation of the highest priority task  $\tau_1$  coincides with the completion of job  $\iota_{3,1}$  of task  $\tau_3$  at time  $t = 0$ . Furthermore, task  $\tau_2$  is activated a small time  $\Delta > 0$  after the start time of  $\iota_{3,1}$ .

### 3.2. An introductory example

Based on the notion of optimal instant for FPPS, it could be tentative to think that an optimal instant for FPTs may occur when the completion of a job of a task  $\tau_i$  coincides with the simultaneous activation of all higher priority tasks that can preempt  $\tau_i$ . These are the tasks in  $\{\tau_h | \pi_h > \theta_i\}$ . In this section, we introduce an example that refutes this intuition of optimal instant for FPTs.

Consider the set  $\mathcal{T}_1$  of three tasks with characteristics as described in Table 1. The best-case response time values in this table were found using brute force. This is, by considering the phasings  $\phi_1 \in [0, 35)$  and  $\phi_2 \in [0, 50)$  with a granularity of one time unit in the simulations, while fixing the phasing of  $\tau_3$  to  $\phi_3 = 0$ .

TABLE 1. TASK CHARACTERISTICS OF  $\mathcal{T}_1$ .

task	$T_i$	$WD_i$	$WC_i = BC_i$	$\pi_i$	$\theta_i$	$WR_i$	$BR_i$
$\tau_1$	35	20	10	3	3	10	10
$\tau_2$	50	65	20	2	2	62	20
$\tau_3$	70	70	22	1	2	66	32

The least common multiple of the periods is 350 and  $U^{\mathcal{T}_1} = 1$ .

Figure 2 shows a timeline for task-set  $\mathcal{T}_1$  with the configuration that we aim to refute for the last job of task  $\tau_3$ . As can be seen, the completion of this job coincides with an activation of the highest priority task  $\tau_1$  at time  $t = 0$ . Note that task  $\tau_1$  can preempt task  $\tau_3$  because  $\pi_1 > \theta_3$ . On the other hand, task  $\tau_2$  cannot preempt task  $\tau_3$  because  $\pi_2 = \theta_3$ . For the latter reason, the activation of task  $\tau_2$  is set to a sufficiently small amount of time  $\Delta > 0$  after the start time of the last job of task  $\tau_3$ , therefore, allowing more time before  $t = 0$  for the execution of task  $\tau_3$ . This is in accordance with the notion of  $\Delta$ -optimal instant.

As can be seen from Figure 2, the last job of task  $\tau_3$  that is experiencing the new optimal instant has a response time of  $R_{3,1} = 36$ . However, the best-case response time of task  $\tau_3$  is  $BR_3 = 32$ , and it is assumed by the job activated at time  $t = -176$ . Note also that the response time of the job experiencing the optimal instant cannot be reduced by pushing the activations of task  $\tau_3$  because the job activated at time  $t = -176$  already can start upon activation. Hence, pushing the activations of task  $\tau_3$  to a later moment in time would modify the schedule.

Based on the previous observations, we formulate the following fact.

**Fact 1.** *In FPTs, a job  $k$  of task  $\tau_i$  does not necessarily assume the best-case response time when the activation of all higher priority tasks that can preempt task  $\tau_i$  coincides with the completion of job  $k$ .*

From Figure 2, we can also observe that, in order to calculate the response time of job  $\iota_{3,1}$ , we have to look at previous jobs in the level-3 active period to see whether they induce some delay in the start time of  $\iota_{3,1}$ . This is similar to the best-case response time analysis for FPPS in [8]. An important property for this analysis, however, is that under FPPS the last job in a level- $i$  active period always experiences the shortest response time of all jobs in that level- $i$  active period. This property clearly does not hold for FPTs as shown in Figure 2. As can be seen, the first job in the level-3 active period is the one with the shortest response time. Therefore, we introduce the following fact.

**Fact 2.** *In FPTs, the shortest response time of a task  $\tau_i$  in a level- $i$  active period is not necessarily assumed by the last job in that level- $i$  active period.*

### 4. Best-case response time for FPNS

Since FPNS is a special case of FPTs when thresholds of tasks are set to the highest priority, we first investigate the best-case response time for non-preemptive tasks as an intermediate step towards the best-case analysis for FPTs.

**Lemma 1.** *Let  $\tau_i$  be a non-preemptive task, i.e.  $\theta_i = \pi_i$ , and let  $U^{\mathcal{T}} < 1$  or  $U^{\mathcal{T}} = 1$  and the lcm of the periods exists. The best-case response time of  $\tau_i$  is always equal to its best-case computation time, i.e.  $BR_i = BC_i$ .*

*Proof.* Given a task-set  $\mathcal{T}$  and a non-preemptive task  $\tau_i \in \mathcal{T}$ , assume that all jobs of  $\tau_i$  have a computation time equal to  $BC_i$ . Since  $\tau_i$  is non-preemptive, the hold time of all its jobs is simply equal to  $BC_i$ . Recall that the response time of a job  $k$  of  $\tau_i$  is given by  $R_{i,k} = S_{i,k} + H_{i,k}$ . Therefore, in order to prove the lemma, it is sufficient to show that it is always possible to schedule task  $\tau_i$  in such a way that the relative start time of a job  $k$  of  $\tau_i$  is zero, i.e.  $S_{i,k} = 0$ . We divide this proof in two cases:

$\{Case U^T < 1\}$ . Given an arbitrary schedule for  $\mathcal{T}$ , let  $[t_s, t_e]$  with  $t_e > t_s$  be an idle interval in such a schedule, where an idle interval is a time interval in which there is no pending work. Note that it is always possible to find an idle interval because  $U^T < 1$ . Furthermore, let  $\iota_{i,k}$  be the first job of  $\tau_i$  activated at or after time  $t_e$ , i.e.  $t_e \leq a_{i,k}$ . We now show that, by pushing the activation of  $\iota_{i,k}$  to occur in the interval  $[t_s, t_e]$ , the relative start time of  $\iota_{i,k}$  becomes zero.

First observe that there is no other job of  $\tau_i$  between the idle interval  $[t_s, t_e]$  and  $\iota_{i,k}$ . Hence, after pushing the activations of  $\tau_i$  till  $a_{i,k}$  occurs in  $[t_s, t_e]$ , job  $\iota_{i,k}$  does not experience blocking by a previous job. Furthermore, since  $[t_s, t_e]$  was originally idle, job  $\iota_{i,k}$  can start upon activation, leading to  $S_{i,k} = 0$ .

$\{Case U^T = 1 \text{ and the lcm of the periods exists}\}$ . First recall that the hyperperiod is defined as the length of the shortest time interval in which the schedule repeats itself after an initial start-up. For a set of periodic tasks, Leung and Merrill proved in [12] that the hyperperiod  $H$  can be calculated as the least common multiple of the periods, i.e.  $H = lcm(T_1, \dots, T_n)$  for  $n$  periodic tasks.

Let  $n_i = H/T_i$  be the number of jobs of  $\tau_i$  in every hyperperiod and assume that no job of  $\tau_i$  starts upon activation. Hence, for each interval  $[a_{i,k'}, s_{i,k'}]$  with  $1 \leq k' \leq n_i$ , there are only higher priority tasks of  $\tau_i$  and/or at most one blocking lower priority task executing within such an interval. As an example of this situation, consider the timeline depicted in Figure 3 for task-set  $\mathcal{T}_2$  described in Table 2. As can be seen, no job of  $\tau_2$  can start upon activation.

TABLE 2. TASK CHARACTERISTICS OF  $\mathcal{T}_2$ .

task	$T_i$	$WD_i$	$C_i$	$\pi_i$	$\theta_i$
$\tau_1$	4	4	2	2	2
$\tau_2$	6	6	3	1	2

The hyperperiod is  $H = 12$  and  $U^{\mathcal{T}_2} = 1$ .




Figure 3. A timeline for  $\mathcal{T}_2$ . The activations of the tasks are repeated in every interval with length  $H = 12$ .

Now let  $\iota_{i,k}$  be the job of  $\tau_i$  with the shortest relative start time among all jobs of  $\tau_i$  in an hyperperiod. Since no job of  $\tau_i$  starts upon activation, we can push the activation of all jobs of  $\tau_i$  to a later moment in time, without modifying the schedule, till job  $\iota_{i,k}$  starts upon activation. After pushing the activations of  $\tau_i$ , it holds that  $S_{i,k} = 0$  therefore concluding the proof. Figure 4 shows an example where the schedule shown in Figure 3 is preserved after pushing the activations of  $\tau_2$  till its second job starts upon activation.  $\square$




Figure 4. A timeline for  $\mathcal{T}_2$ . After pushing the activations of  $\tau_2$ , the second job can immediately start.

## 5. Conclusion

In this paper, we showed that the best-case response time of a non-preemptive task is equal to its execution time when scheduled under FPTs or FPNS. In addition, an initial and intuitive notion of optimal instant for FPTs was refuted. Furthermore, we showed that a crucial property used to develop the best-case response time analysis for FPPS does not longer holds for FPTs. Therefore, making the best-case analysis for FPTs most likely not a trivial extension of the current analysis for FPPS. An exact best-case response time analysis for FPTs is currently under study.

## References

- [1] Express Logic Inc. ThreadX Technical Features, v5.1. [Online]. Available: <http://rtos.com/products/threadx>, June 2010.
- [2] AUTOSAR – Specification of Operating System, Release 4.1. 2010. [Online], Available: <http://www.autosar.org/>
- [3] OSEK group, OSEK/VDX Operating System. February 2005. [Online], Available: <http://portal.osek-vdx.org/files/pdf/specs/os223.pdf>.
- [4] Y. Wang and M. Saksena. Scheduling fixed-priority tasks with preemption threshold. In Proc. 6th International Conf. Real-Time Computing Systems and Applications. (RTCSA), pp. 328-335, Dec. 1999.
- [5] J. Regehr. Scheduling tasks with mixed preemption relations for robustness to timing faults. In Proc. 23rd IEEE Real-Time Systems Symposium (RTSS), pp. 315-326. December 2002.
- [6] U. Keskin, R.J. Bril, and J.J. Lukkien. Exact response-time analysis for fixed-priority preemption-threshold scheduling. In IEEE Conf. Emerging Technologies and Factory Automation (ETFA), Sep. 2010, pp. 14.
- [7] O. Redell and M. Sanfridson. Exact best-case response time analysis of fixed priority scheduled tasks. In Proc. 14th Euromicro Conference on Real-Time Systems (ECRTS), pp. 165-172. June 2002.
- [8] R.J. Bril, J.J. Lukkien, and R.H. Mak. Best-case response times and jitter analysis of real-time tasks with arbitrary deadlines. In Proc. 21st International Conference on Real-Time Networks and Systems (RTNS), ACM, pp. 193-202, October 2013.
- [9] R.I. Davis, A. Burns, R.J. Bril, J.J. Lukkien. Controller area network (CAN) schedulability analysis: refuted, revisited and revised. In Real-Time Systems journal, vol. 35, num. 3, pp. 239-272, 2007.
- [10] R.J. Bril, J.J. Lukkien, W.F.J. Verhaegh. Worst-case response time analysis of real-time tasks under fixed-priority scheduling with deferred preemption. In Real-Time Systems journal, vol. 42, num. 1-3, pp. 63-119. August 2009.
- [11] R.J. Bril, W.F.J. Verhaegh. Towards best-case response times of real-time tasks under fixed-priority scheduling with deferred preemption. In Proc. Work-in-Progress (WiP) session of the 17th Euromicro Conf. on Real-Time Systems (ECRTS), pp. 17-20, July 2005.
- [12] J. Leung, R. Merrill. A note on the preemptive scheduling of periodic, real-time tasks. Information Processing Letters, 18:115-118, 1980.